

Memory-mapped Files

Exceeding NIO Low-latency Capabilities
Memory-mapped Files

Intro

Stream I/O has been improved by leveraging NIO; however, a Memory-mapped File's data access is dramatically faster (by one to two orders of magnitude).

Although Memory-mapped Write operations seem to utilize an instance of `FileOutputStream`, in actuality all of a Memory-mapped File's output uses an instance of `RandomAccessFile`.

The demo application being discussed in this document ([MappedIO](#)) is running in Java 11, it has a side benefit of proving the Java 11 JVM's viability.

Because Memory-mapped Files are related to the Java Memory Model and to GC, these topics are covered in some detail, prior to their application in the demo app.

The [MappedIO](#) demo application is deceptively straightforward, so the underlying concepts will foster a better understanding of the application's findings.

Application Overview

The `start()` *Template Method* in this class iterates through various implementations of the `conductTest()` method, providing a test battery, as defined in the several anonymous inner classes.

Each `TestModule` instance performs one kind of I/O test. The formulation of each of the `conductTest()` methods measures the relative performance of the two I/O approaches. There is a test for both read and write modalities:

Sample output

Stream Write:	36
Mapped Write:	63
Stream Read:	81
Mapped Read:	11
Stream Read/Write:	1673
Mapped Read/Write:	18

Note that the implementation of each `conductTest()` method includes *Wall Time* (this includes the time expended for the initialization of the various I/O objects).

After accounting for scaffolding activities – Memory-mapped Files can be expensive to initialize – the overall throughput gain realized by adopting Memory-mapped Files as compared to standard NIO Streams is significant.

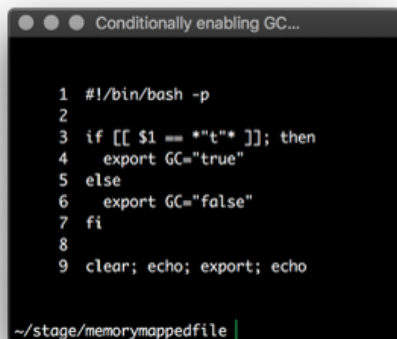
There is also some memory overhead to consider when looking at the effects on the heap in the aftermath of any GC cycle. For instance, you must factor-in *Wired/Resident Memory*, which is reserved for kernel code and for Operating System data structures.

Garbage Collection

Because Memory-mapped Files are intimately related to a JVM's Heap, this demo application allows you to conditionally monitor GC cycles as they occur.

Allowing GC cycles to be reported (they are always running) also messes with the presentation (the console-based representation) of the application's test results. Additionally, printing to console (the `System.out` stream) incurs a time expense/distortion.

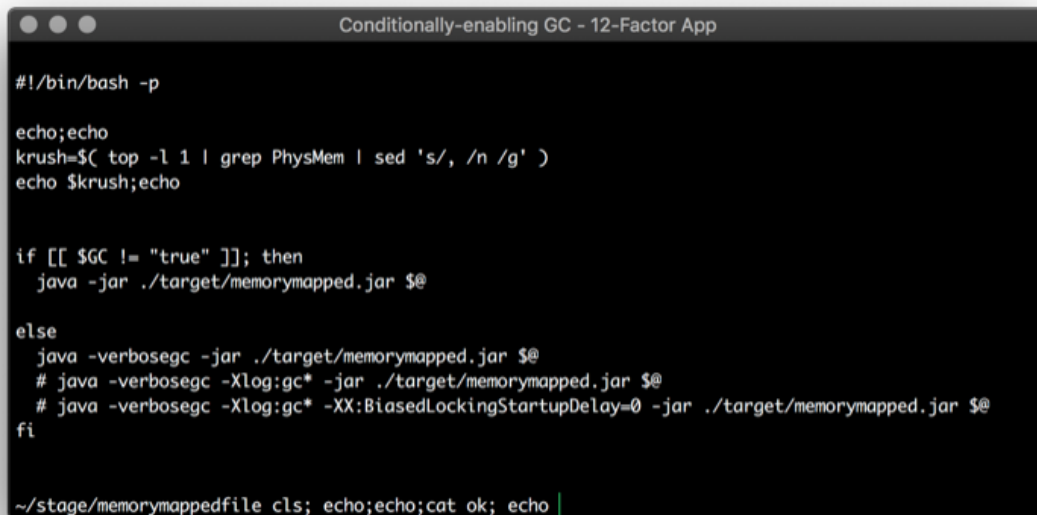
For these reasons, the user can elect to publish the GC cycles, via a CLI switch for the Bash script named `gc`, shown next:



```
1 #!/bin/bash -p
2
3 if [[ $1 == "t" ]]; then
4     export GC="true"
5 else
6     export GC="false"
7 fi
8
9 clear; echo; export; echo

~/stage/memorymappedfile
```

The (12-Factor aware) Bash script named `ok`, shown next, recognizes the value stored in the `$GC` environment variable, whose value is toggled by the `gc` Bash script that was shown above.



```
#!/bin/bash -p

echo;echo
krush=$( top -l 1 | grep PhysMem | sed 's/, /n /g' )
echo $krush;echo

if [[ $GC != "true" ]]; then
    java -jar ./target/memorymapped.jar $@
else
    java -verbosegc -jar ./target/memorymapped.jar $@
    # java -verbosegc -Xlog:gc* -jar ./target/memorymapped.jar $@
    # java -verbosegc -Xlog:gc* -XX:BiasedLockingStartupDelay=0 -jar ./target/memorymapped.jar $@
fi

~/stage/memorymappedfile cls; echo;echo;cat ok; echo
```

Garbage Collection – a deep-dive

As part of this section – aside from discussing GC in the [MappedIO](#) demo application – the topic of Java 11-12 JVMs’ support for [Compressed Oops](#), the [Virt-zero address](#), and [Colored 64-bit Pointers](#) (an unfortunate naming choice) are explored, because the demo is running JVM 11.

Those concepts relate to the Java Memory Model and to GC (the ZGC garbage collector).

1/3 The App and GC

The [MappedIO](#) demo application conditionally allows reporting of (semi) verbose GC activities to be presented alongside the application-oriented readout. Because it is delivered as another thread of execution, the `gc` statements “crash” into the `System` class’ `out/err` Streams (console output).

Below is a representation of a typical (console) gc cycle, *running under the Java 11 JVM*.

The console representation, by the way, depicts a minor GC, as specified by the following canonical CLI switch:

```
-verbosegc
```

```
0.008s][info ][gc,heap] Heap region size: 1M
0.010s][info ][gc    ] Using G1
0.010s][info ][gc,heap,coops] Heap address: 0x0000000700000000, size: 4096 MB, Compressed Oops mode: Zero based, Oop shift amount: 3
```

Breakdown ¶¶

[0.008s]	[info]	[gc,heap]	Heap region size: 1M
[0.010s]	[info]	[gc]	Using G1
[0.010s]	[info]	[gc,heap,coops]	Heap address: 0x0000000700000000, size: 4096 MB, Compressed Oops mode: Zero based, Oop shift amount: 3

2/3 Virt-zero Address

If the OS supports it (post Java 6 does), heap allocation can begin at a virt-zero address – then, Compressed Oop Mode (COop) can be used. COop eliminates the need to add to the base heap address when decoding a 64-bit pointer.

With COop (on a 64-bit architecture) decoding a 64-bit pointer starts with a 32-bit *object* offset, and uses a (3) byte offset as an alignment (the *Oop shift amount* shown in the 1/3 **The App and GC** section, above).

Since a 64-bit pointer uses only 42 bits – and the 3-byte offset, the *Oop shift amount* – the extra bits are used for GC metadata (in a version 11/12 JVM).

Because they're object offsets rather than byte offsets, a 64-bit pointer can address up to four billion *objects* (not bytes) – this equates to a heap size of up to about 32 gigabytes.

Using the COop scheme within an ILP64 architecture, the size of 64-bit pointers are *comparable* to pointers used in ILP32 mode.

To use COop, pointers must be scaled by a factor of 8 and **added to the Java heap base address** to find the object to which they refer.

The computational savings – of decoding/deriving an effective 64-bit native address (pointer) from a 32-bit COop (compressed) address requires the OS being capable of leveraging a **virt-zero address**:

```
[0.008s] [info] [gc,heap] Heap region size: 1M
[0.010s] [info] [gc] Using G1
[0.010s] [info] [gc,heap,coops] Heap address: 0x0000000700000000,
size: 4096 MB,
Compressed Oop mode: Zero based,
Oop shift amount: 3
```

3/3 Colored 64-bit Pointers

ZGC is the latest (low-latency, 10-ms) garbage collector. It trades disk usage for speed. It uses 64-bit colored pointers. It supports varying sized JVMs. *It is experimental.*

Recall that a 64-bit pointer has extra, unused bits (only 42 bits are used).

Hence, the 11-12 versions of the JVM appropriate the extra bits to provide the concept of colored pointers (an unfortunate name). Colored pointers utilize the extra memory (bits) to model the GC-related meta-properties that are listed (in high-order, to low-order bits) below:

```
finalizable
remapped
marked 1
marked 0
```

Marked bits are used for the swapping that occurs between the two Survivor GC spaces

Note about Javadoc Generation

When running w/ Java 11, the [Javadoc Tool](#) no longer supports links that toggle the familiar [frames/noframes](#) feature.

Instead, it creates a [Search](#) box (which appears in the upper-right corner of all generated Javadoc pages).

ref: <https://bugs.openjdk.java.net/browse/JDK-8215599>

Below is a “garden” example of the CLI way to generate Javadoc (Lines [10-12](#), [14-17](#)).

A Portable CLI Javadoc Generation Script

```
1  #!/bin/bash -p
2
3  # Detect the source directory of a Bash script from within the script itself
4  # https://bit.ly/2NokAsM
5
6  # Running Javadoc
7  # https://bit.ly/31XVxSr
8
9
10 ORIGIN="${ cd "$( dirname "${BASH_SOURCE[0]}" )" >/dev/null 2>&1 && pwd }"
11
12 cd $ORIGIN/src/main/java
13
14 javadoc -private -d $ORIGIN/html \
15 -windowtitle 'Memory-mapped Files, Demo' \
16 -doctitle 'Memory-mapped Files, Demo' \
17 org.ascension.util.memorymappedfile.demo
18
19 # open $ORIGIN/html/org/ascension/util/memorymappedfile/demo/MappedIO.html
20
21 cd $ORIGIN
22
23
24 # NB
25 # Under Java 11, the Javadoc tool no longer supports links that
26 # toggle 'frames/noframes'; instead, it supports a 'Search' box,
27 # located in the upper-right corner to replace the links.
28 #
29 # ref: https://bugs.openjdk.java.net/browse/JDK-8215599

~/stage/memorymappedfile |
```

Application Details

The [MappedIO](#) application contrasts two approaches to reading and writing data (files). One approach is streaming using standard NIO; the other approach leverages Memory-mapped Files.

This app shows that the improvement in throughput for Memory-mapped Files over NIO Streams varies from one order of magnitude (for reads) to two orders of magnitude (for a read, then a write).

Besides realizing better throughput, Memory-mapped Files allow your application to access (to page-in) more data than you have RAM.

Additionally – synchronization costs aside – Memory-mapped Files allow multiple threads to share the (paged-in) data. Memory-mapped Files can rapidly access (page-in) other non-resident parts of the file.

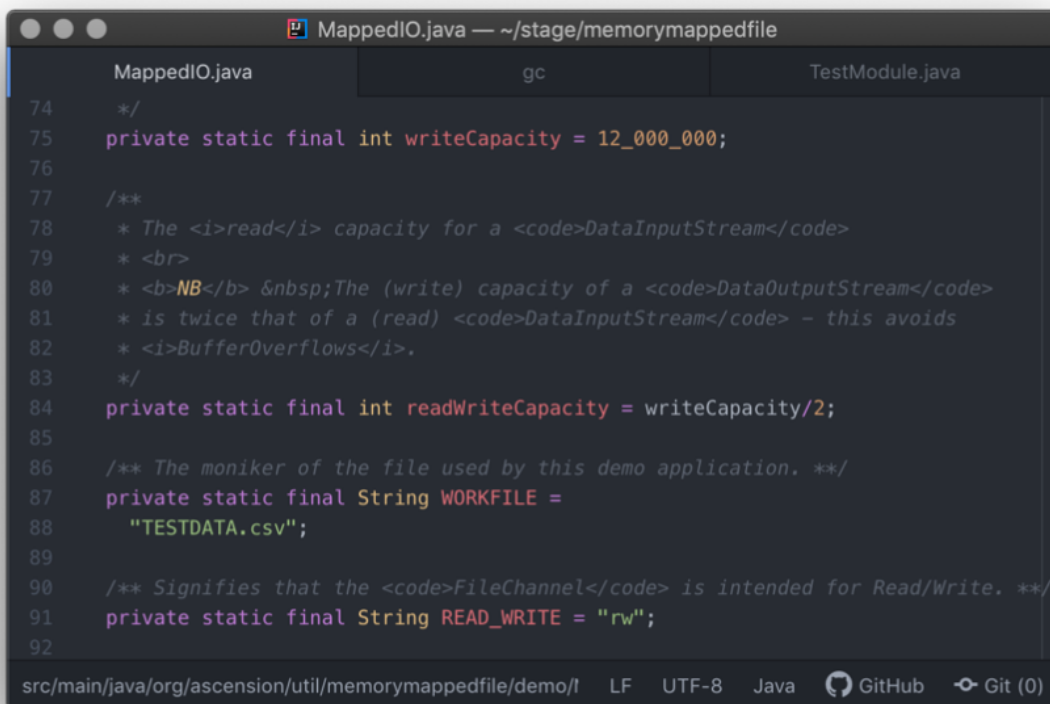
The Memory-mapped arrangement, from a throughput standpoint, is a competitive alternative to using Event-sourcing/Ring Buffers, which requires an Event-Processing architecture topology.

Referring to [Figure One](#) below, ([Lines 75, 84](#)) establish the capacity of the data which is to either be read and/or to be written. These are the three use cases being illustrated by the app.

Note that the write capacity ([Line 75](#)) is twice that of the read (capacity, [Line 84](#)). This provision avoids incurring an `java.nio.BufferOverflowException` at runtime.

The actual value of the write capacity variable is deliberately tagged close to the disk allocation value of the test file. This deliberate sizing level-sets the relative performance of all six (6) tests.

This consideration eliminates any testing bias due to the fact that some tests only write data, some tests only read data (from the test data file), and some test do both.



```
MappedIO.java      gc      TestModule.java
74  */
75  private static final int writeCapacity = 12_000_000;
76
77  /**
78   * The <i>read</i> capacity for a <code>DataInputStream</code>
79   * <br>
80   * <b>NB</b> The (write) capacity of a <code>DataOutputStream</code>
81   * is twice that of a (read) <code>DataInputStream</code> - this avoids
82   * <i>BufferOverflows</i>.
83   */
84  private static final int readWriteCapacity = writeCapacity/2;
85
86  /** The moniker of the file used by this demo application. */
87  private static final String WORKFILE =
88      "TESTDATA.csv";
89
90  /** Signifies that the <code>FileChannel</code> is intended for Read/Write. */
91  private static final String READ_WRITE = "rw";
92
```

src/main/java/org/ascension/util/memorymappedfile/demo/t LF UTF-8 Java GitHub Git (0)

Figure One



The image shows a screenshot of an IDE window titled "MappedIO.java — ~/stage/memorymappedfile". The window contains two tabs: "MappedIO.java" and "TestModule.java". The "MappedIO.java" tab is active, displaying the following Java code:

```
111
112  /** Battery of tests. */
113  private static TestModule[] testBattery = {
114
115      new TestModule(
116          "\n\n  [0] Stream Write\n          Capacity " +
117          writeCapacity/1_000_000 + " MiB") {
118
119          public void conductTest() {
120
121              try (
122
123                  DataOutputStream dataOutputStream =
124                      new DataOutputStream(
125                          new BufferedOutputStream(
126                              new FileOutputStream(
127                                  new File(WORKFILE)))) {
128
129                  for (int i = 0; i < writeCapacity; i++) {
130                      dataOutputStream.writeChar(i);
131                  }
132              }
133              catch (FileNotFoundException e) { e.printStackTrace(); }
134              catch (IOException e) { e.printStackTrace(); }
135          }
136      },
137      new TestModule(
138          "\n  [1] Mapped Write\n          Capacity " +
139          writeCapacity/1_000_000 + " MiB") {
140
141          public void conductTest() {
142
143              try (
144
145                  FileChannel fileChannel =
146                      new RandomAccessFile(WORKFILE, READ_WRITE).
147                      getChannel() {
148
149                  CharBuffer charBuffer = fileChannel.map(
150                      FileChannel.MapMode.READ_WRITE, 0, fileChannel.size()).
151                      asCharBuffer();
152
153                  for (int i = 0; i < writeCapacity; i++) {
154                      charBuffer.put(Character.valueOf((char)i));
155                  }
156              }
157              catch (FileNotFoundException e) { e.printStackTrace(); }
158              catch (IOException e) { e.printStackTrace(); }
159          }
160      }
161  }
```

The bottom status bar shows the file path "src/main/java/org/ascension/util/memorymappedfile/demo", encoding "UTF-8", language "Java", and icons for "GitHub" and "Git (0)".

Figure Two

Line 113 in the screen shot above depicts a collection – a `static` Object array of type:

```
org.ascension.util.memorymappedfile.demo.TestModule.java
```

Each of the successive `TestModule` instances are required to implement the `abstract conductTest()` method. The `conductTest()` method is where each instance fulfills their respective test objectives.

The `TestModule` instances are launched inside a `for (...)` loop, within the `MappedIO` class.

On lines 121-7, 143-7 the canonical `try-with-resources` feature instantiates the particular data structure being utilized for the specific test.

The first code block (Lines 121-7) instantiates a `DataOutputStream`.

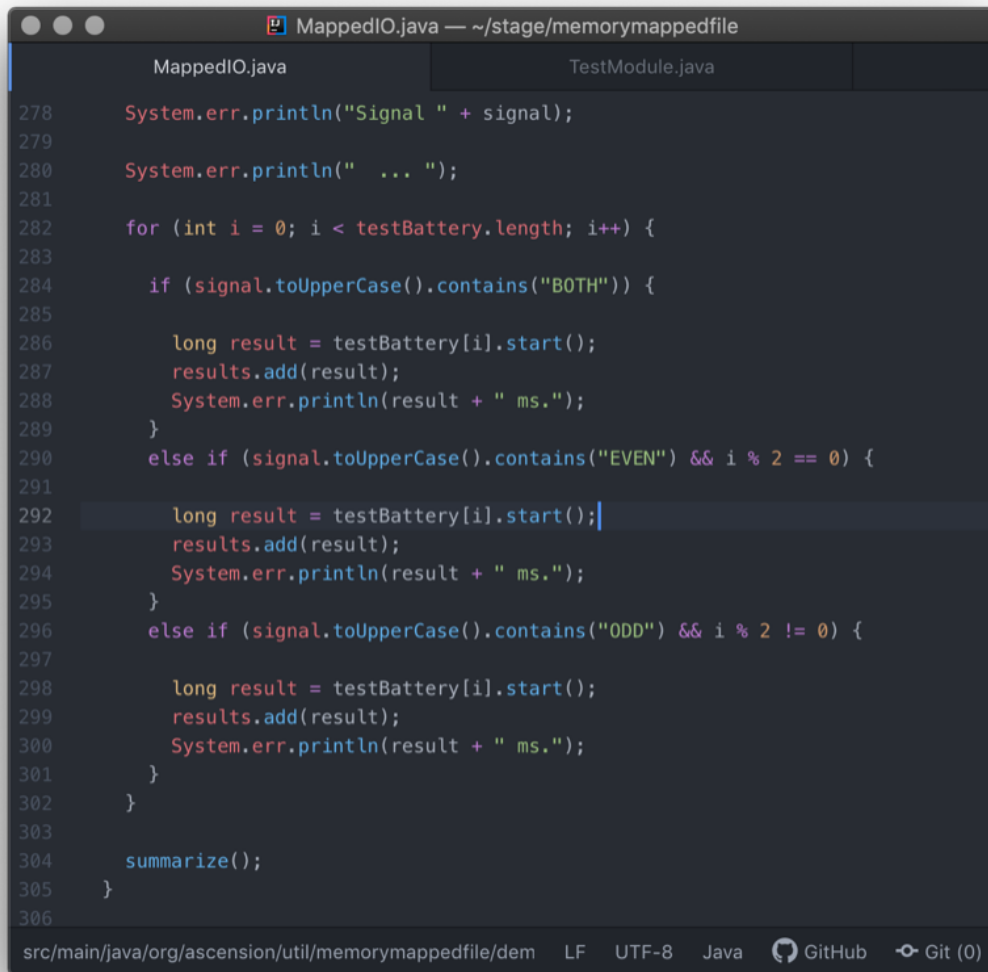
The second code block (Lines 143-7) wraps an instance of (NIO) `java.io.RandomAccessFile` with an instance of `java.nio.channels.FileChannel`.

These two test modules perform “writes” (with identical capacity) – later, the values of their respective throughput characteristics are compared/reported.

This paring of test modules is repeated two more times: a. Reads; b. Reads/Writes.

The Read-Write scenario is the most realistic/common use case for streaming NIO, and for Memory-mapped Files|

The Read-Write scenario shows the most divergence in performance between the two approaches – two orders of magnitude.



```
MappedIO.java — ~/stage/memorymappedfile
MappedIO.java      TestModule.java

278     System.err.println("Signal " + signal);
279
280     System.err.println("  ... ");
281
282     for (int i = 0; i < testBattery.length; i++) {
283
284         if (signal.toUpperCase().contains("BOTH")) {
285
286             long result = testBattery[i].start();
287             results.add(result);
288             System.err.println(result + " ms.");
289         }
290         else if (signal.toUpperCase().contains("EVEN") && i % 2 == 0) {
291
292             long result = testBattery[i].start();
293             results.add(result);
294             System.err.println(result + " ms.");
295         }
296         else if (signal.toUpperCase().contains("ODD") && i % 2 != 0) {
297
298             long result = testBattery[i].start();
299             results.add(result);
300             System.err.println(result + " ms.");
301         }
302     }
303
304     summarize();
305 }
306

src/main/java/org/ascension/util/memorymappedfile/dem  LF  UTF-8  Java  GitHub  Git (0)
```

Figure Three

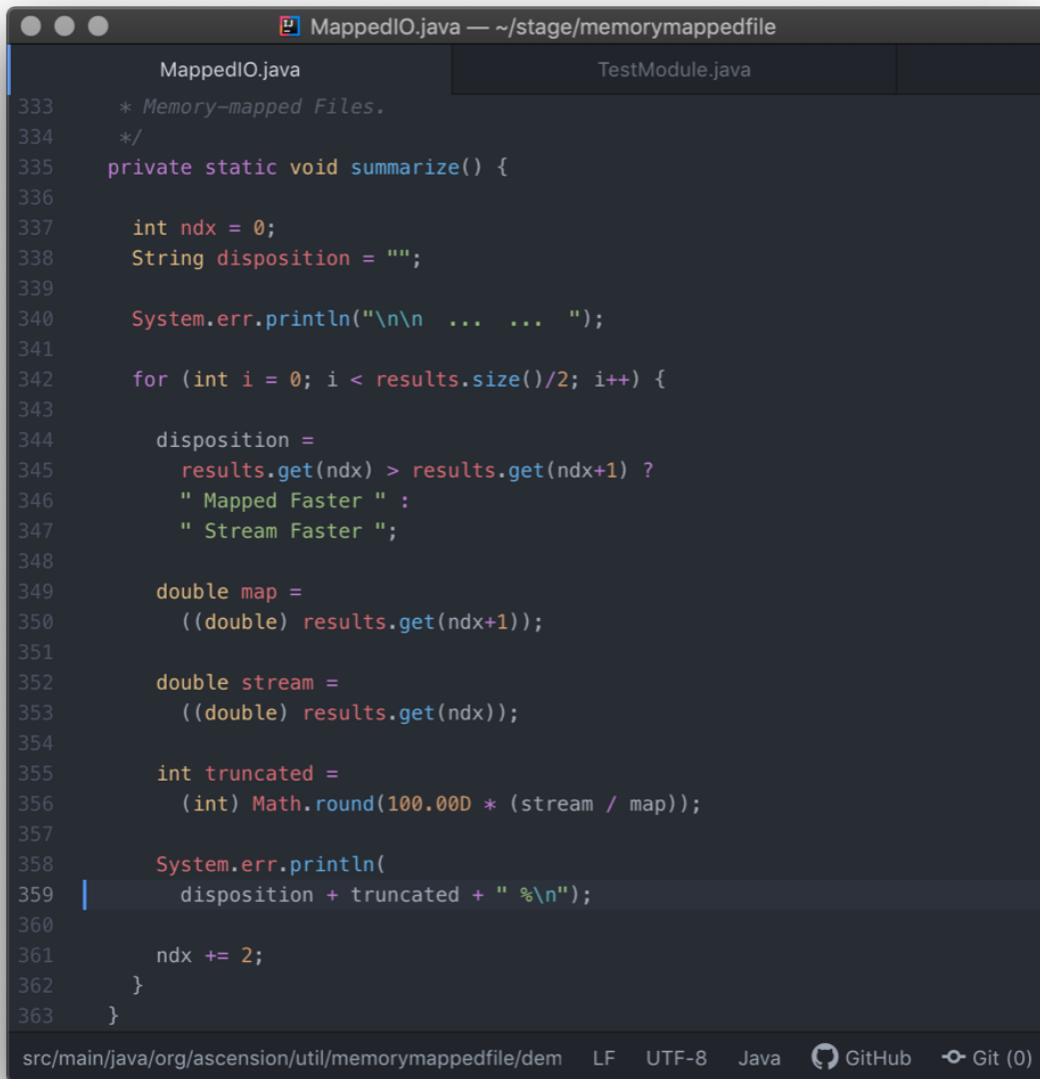
Above, the filtering logic that determines which tests are run (according to CLI input) is shown.

The modulo arithmetic – with a divisor of 2 – determines which **pair** of tests are executed.

If you specify nothing, all six tests are executed.

The `summarize()` method (Line 304) produces a report of the performance differential (as a percentage).

Filtering allows for the elimination of distortion
due to GC which may occur and incur CPU cycles

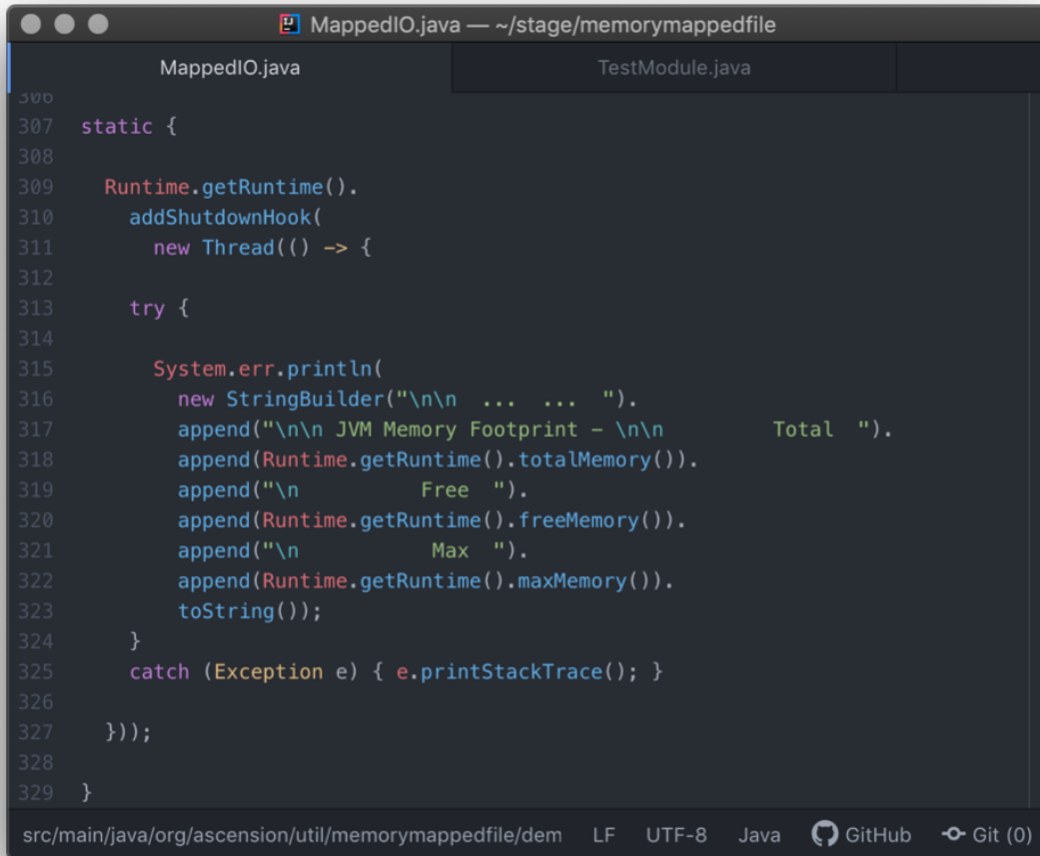


```
333  * Memory-mapped Files.
334  */
335  private static void summarize() {
336
337      int ndx = 0;
338      String disposition = "";
339
340      System.err.println("\n\n ... .. ");
341
342      for (int i = 0; i < results.size()/2; i++) {
343
344          disposition =
345              results.get(ndx) > results.get(ndx+1) ?
346              " Mapped Faster " :
347              " Stream Faster ";
348
349          double map =
350              ((double) results.get(ndx+1));
351
352          double stream =
353              ((double) results.get(ndx));
354
355          int truncated =
356              (int) Math.round(100.00D * (stream / map));
357
358          System.err.println(
359              disposition + truncated + " %\n");
360
361          ndx += 2;
362      }
363  }
```

src/main/java/org/ascension/util/memorymappedfile/dem LF UTF-8 Java GitHub Git (0)

Figure Four

Above, the logic that compiles the comparison of paired test results.



```
MappedIO.java      TestModule.java
300
307 static {
308
309     Runtime.getRuntime().
310         addShutdownHook(
311             new Thread(() -> {
312
313                 try {
314
315                     System.err.println(
316                         new StringBuilder("\n\n ... .. ").
317                             append("\n\n JVM Memory Footprint - \n\n          Total ").
318                             append(Runtime.getRuntime().totalMemory()).
319                             append("\n          Free ").
320                             append(Runtime.getRuntime().freeMemory()).
321                             append("\n          Max ").
322                             append(Runtime.getRuntime().maxMemory()).
323                             toString());
324                 }
325                 catch (Exception e) { e.printStackTrace(); }
326
327             }));
328
329     }
```

src/main/java/org/ascension/util/memorymappedfile/dem LF UTF-8 Java GitHub Git (0)

Figure Five

This screen shot shows a canonical Java `ShutdownHook`, containing logic that reports the JVM's memory footprint at the end of the application's execution.

Application, in-flights

```
memorymappedfile
[INFO] Finished at: 2019-11-07T15:24:12-06:00
[INFO] -----

PhysMem: 15G used (2624M wired)n 1095M unused.

Signal BOTH
...

[0] Stream Write
    Capacity 12 MiB 256 ms.

[1] Mapped Write
    Capacity 12 MiB 76 ms.

[2] Stream Read
    Capacity 6 MiB 248 ms.

[3] Mapped Read
    Capacity 6 MiB 10 ms.

[4] Stream Read/Write
    Capacity 6 MiB 39543 ms.

[5] Mapped Read/Write
    Capacity 6 MiB 127 ms.

... ..
Mapped Faster 337 %

Mapped Faster 2480 %

Mapped Faster 31136 %

... ..

JVM Memory Footprint -
    Total 298319872
    Free 288025184
    Max 3817865216
```

Outros

The treatment of the Memory-mapped Files (using the [MappedIO](#) app as an example) covers GC topics and discusses salient parts of the Java Language's Memory-model.

The best part is that, as Java evolves, developers benefit from the underlying improvements with no attendant refactoring costs to the code; however, with a deep understanding of improvements, they can design more performant applications.

Appendix

<https://mechanical-sympathy.blogspot.com/2011/11/biased-locking-osr-and-benchmarking-fun.html>

<https://docs.oracle.com/javase/7/docs/technotes/guides/vm/performance-enhancements-7.html#compressedOop>

<https://www.baeldung.com/jvm-compressed-oops>

<https://www.baeldung.com/java-verbose-gc>

<https://codeahoy.com/2017/08/06/basics-of-java-garbage-collection/>

<https://www.opsian.com/blog/javas-new-zgc-is-very-exciting/>

<http://www.diva-portal.org/smash/get/diva2:754541/FULLTEXT01.pdf>

https://subscription.packtpub.com/book/application_development/9781789133271/8/ch08lv1sec45/colored-pointers